

ASSIGNMENT NO:1

```
In [ ]: NAME: Habibsaeed Mukebil  
ROLLNO:14225
```

```
In [30]: import pandas as pd  
import numpy as np  
import seaborn as sns  
import matplotlib.pyplot as plt  
import warnings  
warnings.filterwarnings('ignore')
```

```
In [31]: df=pd.read_csv("uber.csv")  
df
```

Out[31]:

	Unnamed: 0	key	fare_amount	pickup_datetime	pickup_longitude	p
0	24238194	2015-05-07 19:52:06.0000003	7.5	2015-05-07 19:52:06 UTC	-73.999817	
1	27835199	2009-07-17 20:04:56.0000002	7.7	2009-07-17 20:04:56 UTC	-73.994355	
2	44984355	2009-08-24 21:45:00.00000061	12.9	2009-08-24 21:45:00 UTC	-74.005043	
3	25894730	2009-06-26 08:22:21.0000001	5.3	2009-06-26 08:22:21 UTC	-73.976124	
4	17610152	2014-08-28 17:47:00.000000188	16.0	2014-08-28 17:47:00 UTC	-73.925023	
...
199995	42598914	2012-10-28 10:49:00.00000053	3.0	2012-10-28 10:49:00 UTC	-73.987042	
199996	16382965	2014-03-14 01:09:00.0000008	7.5	2014-03-14 01:09:00 UTC	-73.984722	
199997	27804658	2009-06-29 00:42:00.00000078	30.9	2009-06-29 00:42:00 UTC	-73.986017	
199998	20259894	2015-05-20 14:56:25.0000004	14.5	2015-05-20 14:56:25 UTC	-73.997124	
199999	11951496	2010-05-15 04:08:00.00000076	14.1	2010-05-15 04:08:00 UTC	-73.984395	

200000 rows × 9 columns

```
In [32]: df.head()
```

Out[32]:

	Unnamed: 0	key	fare_amount	pickup_datetime	pickup_longitude	pickup_
0	24238194	2015-05-07 19:52:06.0000003	7.5	2015-05-07 19:52:06 UTC	-73.999817	40
1	27835199	2009-07-17 20:04:56.0000002	7.7	2009-07-17 20:04:56 UTC	-73.994355	40
2	44984355	2009-08-24 21:45:00.00000061	12.9	2009-08-24 21:45:00 UTC	-74.005043	40
3	25894730	2009-06-26 08:22:21.0000001	5.3	2009-06-26 08:22:21 UTC	-73.976124	40
4	17610152	2014-08-28 17:47:00.000000188	16.0	2014-08-28 17:47:00 UTC	-73.925023	40

In [33]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200000 entries, 0 to 199999
Data columns (total 9 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Unnamed: 0        200000 non-null   int64  
 1   key              200000 non-null   object  
 2   fare_amount      200000 non-null   float64 
 3   pickup_datetime  200000 non-null   object  
 4   pickup_longitude 200000 non-null   float64 
 5   pickup_latitude  200000 non-null   float64 
 6   dropoff_longitude 199999 non-null   float64 
 7   dropoff_latitude 199999 non-null   float64 
 8   passenger_count  200000 non-null   int64  
dtypes: float64(5), int64(2), object(2)
memory usage: 13.7+ MB
```

In [34]: `df.columns`

```
Out[34]: Index(['Unnamed: 0', 'key', 'fare_amount', 'pickup_datetime',
       'pickup_longitude', 'pickup_latitude', 'dropoff_longitude',
       'dropoff_latitude', 'passenger_count'],
      dtype='object')
```

In [35]: `df = df.drop(['Unnamed: 0', 'key'], axis=1)`

In [36]: `df.head()`

Out[36]:

	fare_amount	pickup_datetime	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude
0	7.5	2015-05-07 19:52:06 UTC	-73.999817	40.738354	-73.999512	
1	7.7	2009-07-17 20:04:56 UTC	-73.994355	40.728225	-73.994710	
2	12.9	2009-08-24 21:45:00 UTC	-74.005043	40.740770	-73.962565	
3	5.3	2009-06-26 08:22:21 UTC	-73.976124	40.790844	-73.965316	
4	16.0	2014-08-28 17:47:00 UTC	-73.925023	40.744085	-73.973082	

◀ ▶

In [37]: `df.shape`

Out[37]: (200000, 7)

In [38]: `df.dtypes`

Out[38]:

fare_amount	float64
pickup_datetime	object
pickup_longitude	float64
pickup_latitude	float64
dropoff_longitude	float64
dropoff_latitude	float64
passenger_count	int64
dtype:	object

In [39]: `df.describe()`

Out[39]:

	fare_amount	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude
count	200000.000000	200000.000000	200000.000000	199999.000000	199999.000000
mean	11.359955	-72.527638	39.935885	-72.525292	39.92389
std	9.901776	11.437787	7.720539	13.117408	6.79482
min	-52.000000	-1340.648410	-74.015515	-3356.666300	-881.98551
25%	6.000000	-73.992065	40.734796	-73.991407	40.73382
50%	8.500000	-73.981823	40.752592	-73.980093	40.75304
75%	12.500000	-73.967154	40.767158	-73.963658	40.76800
max	499.000000	57.418457	1644.421482	1153.572603	872.69762

◀ ▶

In [40]: `df.isnull().sum()`

```
Out[40]: fare_amount      0  
pickup_datetime     0  
pickup_longitude    0  
pickup_latitude     0  
dropoff_longitude   1  
dropoff_latitude    1  
passenger_count     0  
dtype: int64
```

```
In [41]: df['dropoff_latitude'].fillna(value=df['dropoff_latitude'].mean(), inplace=True)
```

```
In [42]: df.isnull().sum()
```

```
Out[42]: fare_amount      0  
pickup_datetime     0  
pickup_longitude    0  
pickup_latitude     0  
dropoff_longitude   1  
dropoff_latitude    0  
passenger_count     0  
dtype: int64
```

```
In [43]: df['dropoff_longitude'].fillna(value=df['dropoff_longitude'].median(), inplace=True)
```

```
In [44]: df.isnull().sum()
```

```
Out[44]: fare_amount      0  
pickup_datetime     0  
pickup_longitude    0  
pickup_latitude     0  
dropoff_longitude   0  
dropoff_latitude    0  
passenger_count     0  
dtype: int64
```

```
In [45]: df.dtypes
```

```
Out[45]: fare_amount      float64  
pickup_datetime     object  
pickup_longitude    float64  
pickup_latitude     float64  
dropoff_longitude   float64  
dropoff_latitude    float64  
passenger_count     int64  
dtype: object
```

```
In [46]: df.pickup_datetime = pd.to_datetime(df.pickup_datetime, errors='coerce')
```

```
In [47]: df.dtypes
```

```
Out[47]: fare_amount           float64
pickup_datetime      datetime64[ns, UTC]
pickup_longitude        float64
pickup_latitude         float64
dropoff_longitude        float64
dropoff_latitude         float64
passenger_count          int64
dtype: object
```

```
In [48]: df=df.assign(hour = df.pickup_datetime.dt.hour,
                     day = df.pickup_datetime.dt.day,
                     month = df.pickup_datetime.dt.month,
                     year = df.pickup_datetime.dt.year,
                     dayofweek = df.pickup_datetime.dt.dayofweek)
```

```
In [49]: df.head
```

```
Out[49]: <bound method NDFrame.head of
p_longitude \
0           7.5 2015-05-07 19:52:06+00:00      -73.999817
1           7.7 2009-07-17 20:04:56+00:00      -73.994355
2          12.9 2009-08-24 21:45:00+00:00      -74.005043
3           5.3 2009-06-26 08:22:21+00:00      -73.976124
4          16.0 2014-08-28 17:47:00+00:00      -73.925023
...
199995      3.0 2012-10-28 10:49:00+00:00      -73.987042
199996      7.5 2014-03-14 01:09:00+00:00      -73.984722
199997     30.9 2009-06-29 00:42:00+00:00      -73.986017
199998     14.5 2015-05-20 14:56:25+00:00      -73.997124
199999     14.1 2010-05-15 04:08:00+00:00      -73.984395

pickup_latitude  dropoff_longitude  dropoff_latitude  passenger_count \
0            40.738354        -73.999512        40.723217           1
1            40.728225        -73.994710        40.750325           1
2            40.740770        -73.962565        40.772647           1
3            40.790844        -73.965316        40.803349           3
4            40.744085        -73.973082        40.761247           5
...
199995      40.739367        -73.986525        40.740297           1
199996      40.736837        -74.006672        40.739620           1
199997      40.756487        -73.858957        40.692588           2
199998      40.725452        -73.983215        40.695415           1
199999      40.720077        -73.985508        40.768793           1

hour  day  month  year  dayofweek
0    19    7      5  2015       3
1    20   17      7  2009       4
2    21   24      8  2009       0
3     8   26      6  2009       4
4    17   28      8  2014       3
...
199995    10   28     10  2012       6
199996     1   14      3  2014       4
199997     0   29      6  2009       0
199998    14   20      5  2015       2
199999     4   15      5  2010       5

[200000 rows x 12 columns]>
```

```
In [50]: df = df.drop('pickup_datetime', axis=1)
```

```
In [51]: df.head
```

```
Out[51]: <bound method NDFrame.head of
de dropoff_longitude \
0           7.5      -73.999817      40.738354      -73.999512
1           7.7      -73.994355      40.728225      -73.994710
2          12.9      -74.005043      40.740770      -73.962565
3           5.3      -73.976124      40.790844      -73.965316
4          16.0      -73.925023      40.744085      -73.973082
...
199995       3.0      -73.987042      40.739367      -73.986525
199996       7.5      -73.984722      40.736837      -74.006672
199997      30.9      -73.986017      40.756487      -73.858957
199998      14.5      -73.997124      40.725452      -73.983215
199999      14.1      -73.984395      40.720077      -73.985508

dropoff_latitude  passenger_count  hour  day  month  year  dayofweek
0           40.723217             1    19    7     5   2015      3
1           40.750325             1    20   17     7   2009      4
2           40.772647             1    21   24     8   2009      0
3           40.803349             3     8   26     6   2009      4
4           40.761247             5    17   28     8   2014      3
...
199995      40.740297             1    10   28    10   2012      6
199996      40.739620             1     1   14     3   2014      4
199997      40.692588             2     0   29     6   2009      0
199998      40.695415             1    14   20     5   2015      2
199999      40.768793             1     4   15     5   2010      5

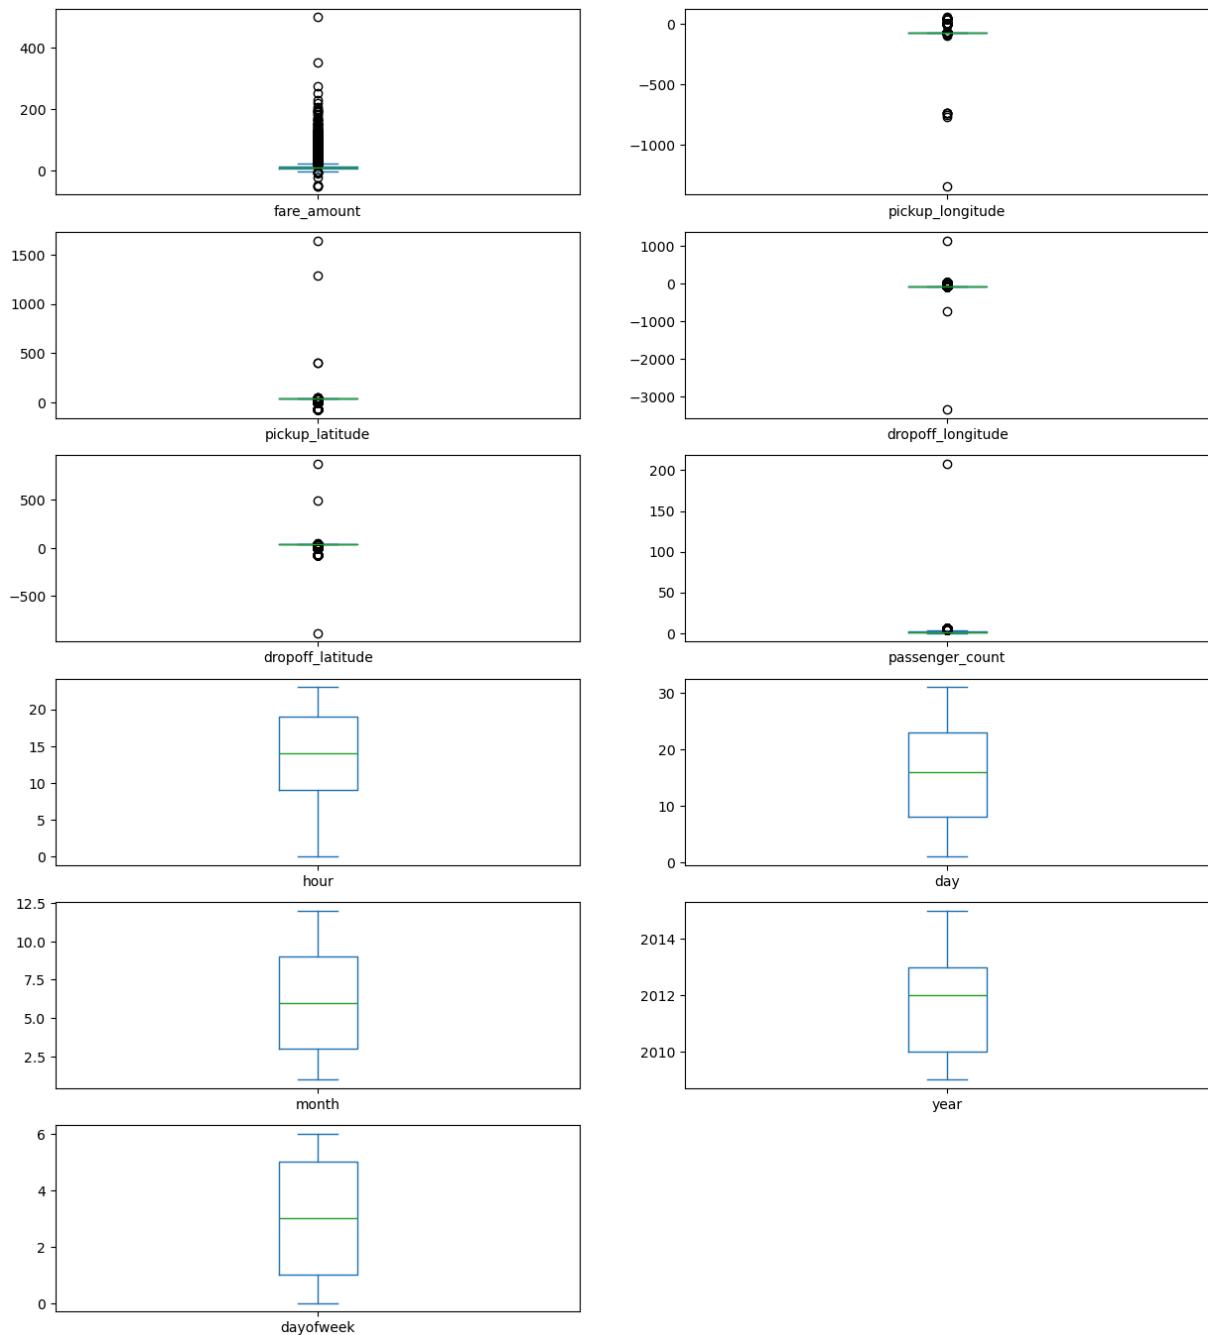
[200000 rows x 11 columns]>
```

```
In [52]: df.dtypes
```

```
Out[52]: fare_amount      float64
pickup_longitude      float64
pickup_latitude       float64
dropoff_longitude      float64
dropoff_latitude       float64
passenger_count        int64
hour                   int32
day                     int32
month                  int32
year                   int32
dayofweek                int32
dtype: object
```

```
In [53]: df.plot(kind="box", subplots = True, layout =(7,2), figsize=(15,20))
```

```
Out[53]: fare_amount      Axes(0.125, 0.786098; 0.352273x0.0939024)
pickup_longitude    Axes(0.547727, 0.786098; 0.352273x0.0939024)
pickup_latitude      Axes(0.125, 0.673415; 0.352273x0.0939024)
dropoff_longitude    Axes(0.547727, 0.673415; 0.352273x0.0939024)
dropoff_latitude     Axes(0.125, 0.560732; 0.352273x0.0939024)
passenger_count      Axes(0.547727, 0.560732; 0.352273x0.0939024)
hour                 Axes(0.125, 0.448049; 0.352273x0.0939024)
day                  Axes(0.547727, 0.448049; 0.352273x0.0939024)
month                Axes(0.125, 0.335366; 0.352273x0.0939024)
year                 Axes(0.547727, 0.335366; 0.352273x0.0939024)
dayofweek             Axes(0.125, 0.222683; 0.352273x0.0939024)
dtype: object
```



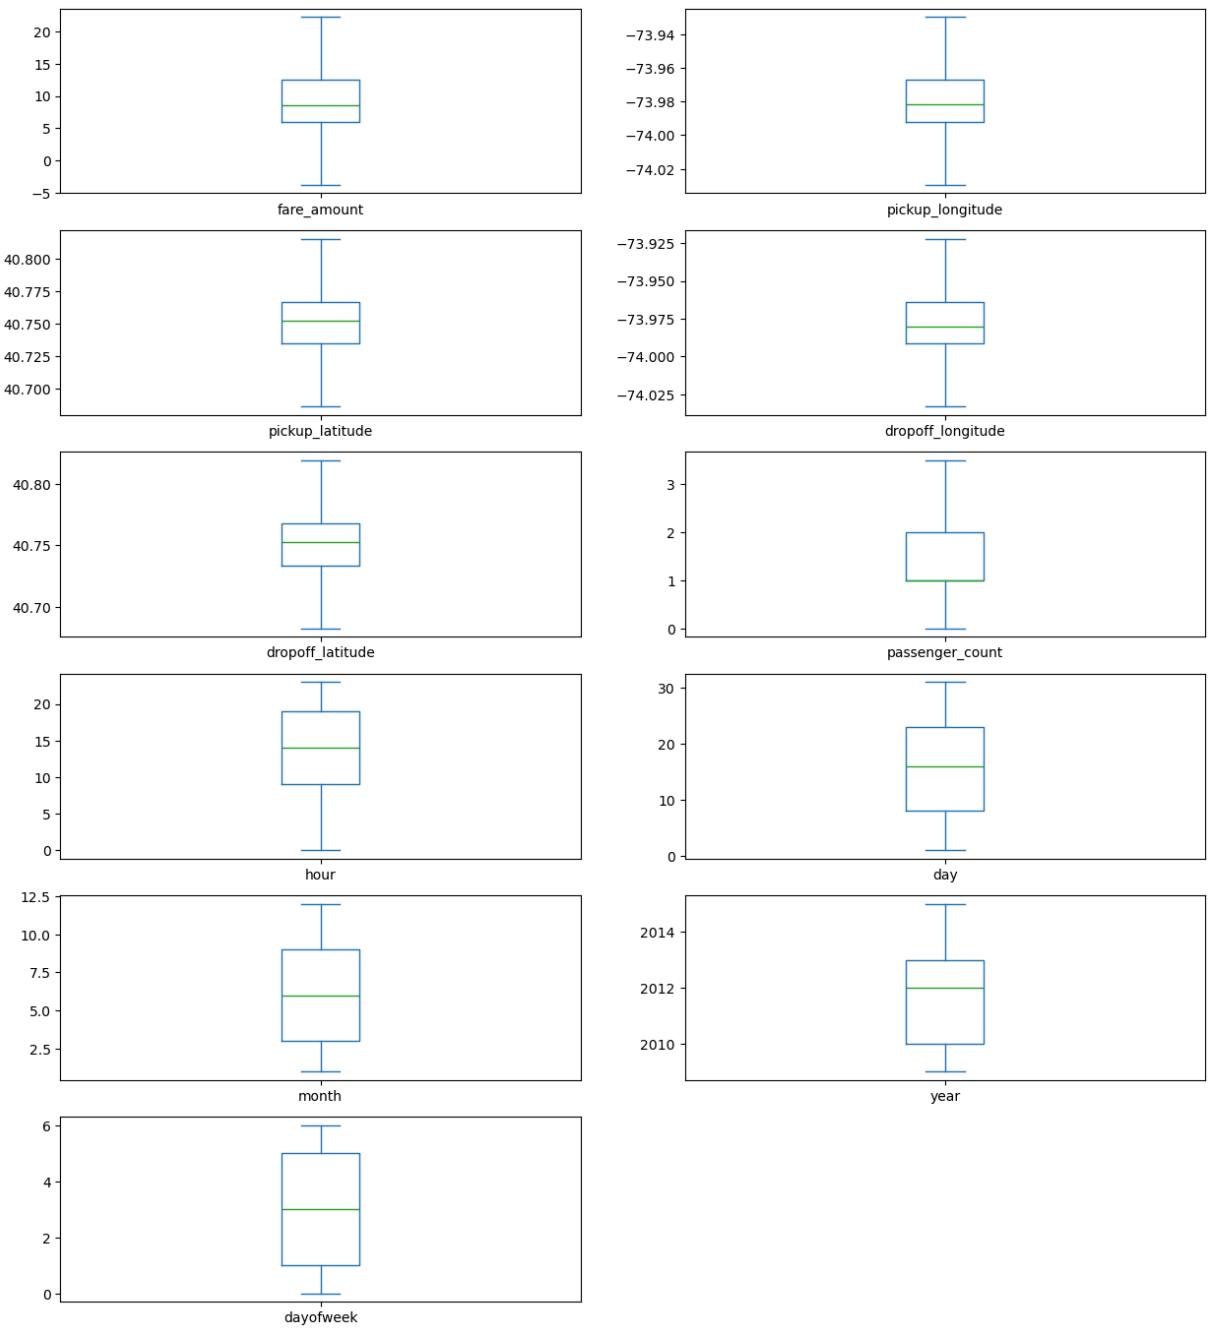
```
In [62]: def remove_outlier(df1, col):
    Q1 = df1[col].quantile(0.25)
    Q3 = df1[col].quantile(0.75)
```

```
IQR= Q3 - Q1
lower_whisker= Q1-1.5*IQR
upper_whisker= Q3+1.5*IQR
df[col] = np.clip(df1[col], lower_whisker, upper_whisker)
return df1
def treat_outliers_all(df1, col_list):
    for c in col_list:
        df1 = remove_outlier(df, c)
    return df1
```

```
In [63]: df = treat_outliers_all(df,df.iloc[:, 0::])
```

```
In [64]: df.plot(kind = "box", subplots = True, layout=(7,2), figsize=(15,20))
```

```
Out[64]: fare_amount      Axes(0.125,0.786098;0.352273x0.0939024)
pickup_longitude     Axes(0.547727,0.786098;0.352273x0.0939024)
pickup_latitude       Axes(0.125,0.673415;0.352273x0.0939024)
dropoff_longitude     Axes(0.547727,0.673415;0.352273x0.0939024)
dropoff_latitude       Axes(0.125,0.560732;0.352273x0.0939024)
passenger_count       Axes(0.547727,0.560732;0.352273x0.0939024)
hour                  Axes(0.125,0.448049;0.352273x0.0939024)
day                   Axes(0.547727,0.448049;0.352273x0.0939024)
month                 Axes(0.125,0.335366;0.352273x0.0939024)
year                  Axes(0.547727,0.335366;0.352273x0.0939024)
dayofweek              Axes(0.125,0.222683;0.352273x0.0939024)
dtype: object
```



```
In [65]: pip install haversine
```

```
Collecting haversine
  Downloading haversine-2.9.0-py2.py3-none-any.whl.metadata (5.8 kB)
Downloaded haversine-2.9.0-py2.py3-none-any.whl (7.7 kB)
Installing collected packages: haversine
Successfully installed haversine-2.9.0
Note: you may need to restart the kernel to use updated packages.
```

```
In [72]: import haversine as hs
travel_dist = []

for pos in range(len(df['pickup_longitude'])):
    long1 = df['pickup_longitude'][pos]
    lati1 = df['pickup_latitude'][pos]
    long2 = df['dropoff_longitude'][pos]
    lati2 = df['dropoff_latitude'][pos]
    travel_dist.append(hs.haversine((long1, lati1), (long2, lati2)))
```

```

lati2 = df['dropoff_latitude'][pos] # <-- corrected from 'dropoff_lloc1'

loc1 = (lati1, long1)
loc2 = (lati2, long2)

c = hs.haversine(loc1, loc2)
travel_dist.append(c)

df['dist_travel_km'] = travel_dist
df.head()

```

Out[72]:

	fare_amount	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude	pas
0	7.5	-73.999817	40.738354	-73.999512	40.723217	
1	7.7	-73.994355	40.728225	-73.994710	40.750325	
2	12.9	-74.005043	40.740770	-73.962565	40.772647	
3	5.3	-73.976124	40.790844	-73.965316	40.803349	
4	16.0	-73.929786	40.744085	-73.973082	40.761247	



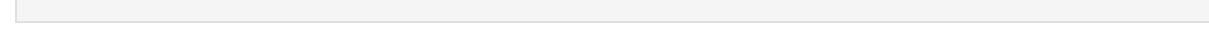
In [73]:

```
df = df.loc[(df.dist_travel_km>=1) | (df.dist_travel_km<=130) ]
print("Remaining obervation: ", df.shape)
```

Remaining obervation: (200000, 12)

In [74]:

```
incorrect_coordinates =df.loc[(df.pickup_latitude>90) | (df.pickup_latitude< -90)| (df.dropoff_latitude>90) | (df.dropoff_latitude< -90) | (df.pickup_longitude>180) | (df.pickup_longitude< -180) | (df.dropoff_longitude>90) | (df.dropoff_longitude< -90) ]
```



In [75]:

```
df.drop(incorrect_coordinates, inplace = True, errors = 'ignore')
```

In [76]:

```
df.head()
```

Out[76]:

	fare_amount	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude	pas
0	7.5	-73.999817	40.738354	-73.999512	40.723217	
1	7.7	-73.994355	40.728225	-73.994710	40.750325	
2	12.9	-74.005043	40.740770	-73.962565	40.772647	
3	5.3	-73.976124	40.790844	-73.965316	40.803349	
4	16.0	-73.929786	40.744085	-73.973082	40.761247	



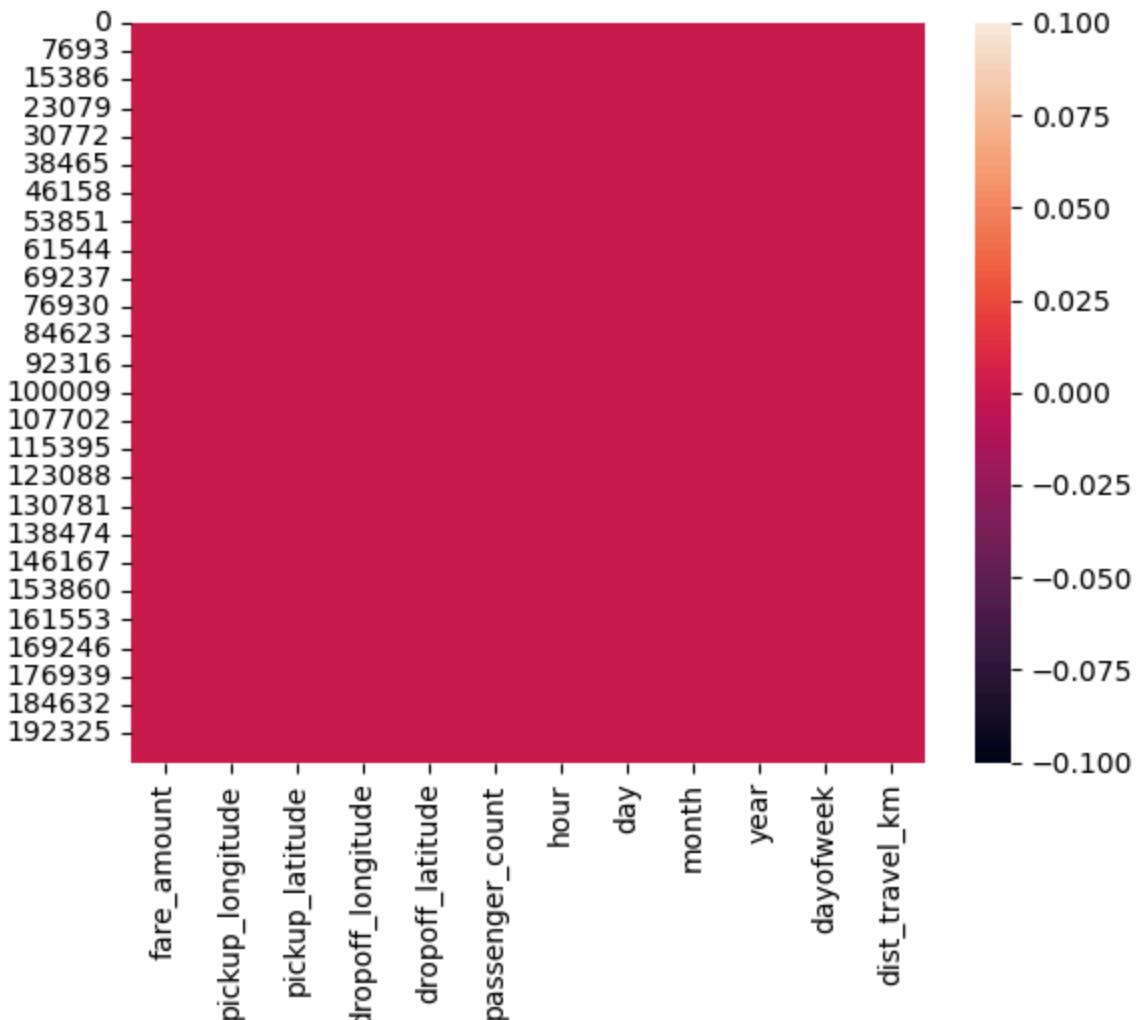
In [77]:

```
df.isnull().sum()
```

```
Out[77]: fare_amount      0  
pickup_longitude     0  
pickup_latitude       0  
dropoff_longitude     0  
dropoff_latitude       0  
passenger_count      0  
hour                  0  
day                   0  
month                 0  
year                  0  
dayofweek              0  
dist_travel_km        0  
dtype: int64
```

```
In [78]: sns.heatmap(df.isnull())
```

```
Out[78]: <Axes: >
```



```
In [79]: corr = df.corr()  
corr
```

Out[79]:

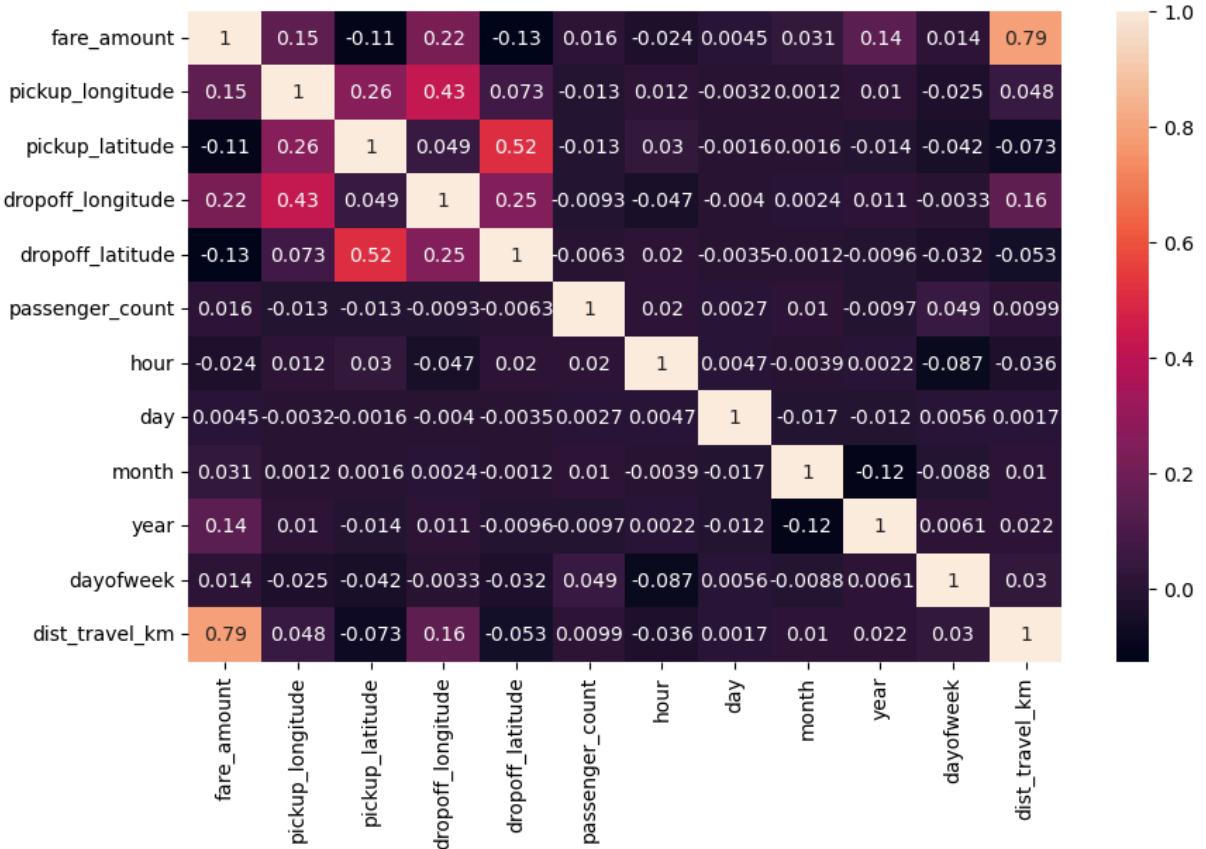
	fare_amount	pickup_longitude	pickup_latitude	dropoff_longitude	drop
fare_amount	1.000000	0.154069	-0.110842	0.218675	
pickup_longitude	0.154069	1.000000	0.259497	0.425619	
pickup_latitude	-0.110842	0.259497	1.000000	0.048889	
dropoff_longitude	0.218675	0.425619	0.048889	1.000000	
dropoff_latitude	-0.125898	0.073290	0.515714	0.245667	
passenger_count	0.015778	-0.013213	-0.012889	-0.009303	
hour	-0.023623	0.011579	0.029681	-0.046558	
day	0.004534	-0.003204	-0.001553	-0.004007	
month	0.030817	0.001169	0.001562	0.002391	
year	0.141277	0.010198	-0.014243	0.011346	
dayofweek	0.013652	-0.024652	-0.042310	-0.003336	
dist_travel_km	0.786385	0.048446	-0.073362	0.155191	



In [80]:

```
fig, axis = plt.subplots(figsize= (10,6))
sns.heatmap(df.corr(), annot = True)
```

Out[80]: <Axes: >



```
In [81]: df.dtypes
```

```
Out[81]: fare_amount      float64
pickup_longitude    float64
pickup_latitude     float64
dropoff_longitude   float64
dropoff_latitude    float64
passenger_count     float64
hour                 int32
day                  int32
month                int32
year                 int32
dayofweek            int32
dist_travel_km       float64
dtype: object
```

```
In [84]: x = df[['pickup_longitude', 'pickup_latitude', 'dropoff_longitude', 'passenger_count']]
```

```
In [85]: y = df['fare_amount']
```

```
In [87]: from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size = 0.33)
```

```
In [89]: from sklearn.linear_model import LinearRegression
regression = LinearRegression()
```

```
In [90]: regression.fit(x_train, y_train)
```

```
Out[90]: ▾ LinearRegression ⓘ ⓘ
LinearRegression()
```

```
In [91]: regression.intercept_
```

```
Out[91]: 3256.513201301484
```

```
In [92]: regression.coef_
```

```
Out[92]: array([ 3.10758620e+01, -1.85910612e+01,  1.27429851e+01,  5.41800674e-02,
                  5.77478080e-03,  3.38249411e-03,  5.89651957e-02,  3.71366494e-01,
                 -3.27278737e-02,  1.86084919e+00])
```

```
In [93]: prediction = regression.predict(x_test)
```

```
In [94]: print(prediction)
```

```
[13.00853888  7.57316671  13.25903141 ... 13.71316448 13.18778424
 11.11547073]
```

```
In [95]: y_test
```

```
Out[95]: 1086    22.25
96074    11.30
104027   12.50
133659    7.50
58255    7.70
...
27046    10.50
178816   14.00
42112    12.50
53060    14.50
83308    8.90
Name: fare_amount, Length: 66000, dtype: float64
```

```
In [97]: from sklearn.metrics import r2_score
r2_score(y_test,prediction)
```

```
Out[97]: 0.6584561379465177
```

```
In [100...]: from sklearn.metrics import mean_squared_error
MSE = mean_squared_error(y_test,prediction)
MSE
```

```
Out[100...]: 10.140972773285915
```

```
In [ ]:
```